

Group Assignment 1: Crafting and learning features

Computer Vision (H02A5a)

The goal of this assignment is to explore more advanced techniques for constructing features that better describe objects of interest and to perform face recognition using these features. This assignment will be delivered in groups of 5 (either composed by you or randomly assigned by your TAs). Start from the provided notebook under *Starter Code* on Toledo (Content > Assignments > Group Assignment 1 > Starter Code). The notebook you submit for grading is the last notebook you submitted in the Kaggle competition prior to the deadline. Make sure that this notebook is well documented. You can submit as many times as you like, so we advise to regularly run the entire notebook on Kaggle. Good luck and have fun!

1 Overview

In the first part of this assignment, you will start from a small dataset and construct/learn features that describe the data. We discriminate two types of features, handcrafted features and features learned from data. Handcrafted features are properties that can be algorithmically extracted from the image (similar to the previous assignment). In this assignment you will experiment with Histogram of Oriented Gradients (HOG) or Scale Invariant Feature Transform (SIFT) feature descriptors. The features learned from data are features that are learned by finding patterns in multiple images. You will explore Principal Component Analysis (PCA) in the context of face detection, so-called eigenfaces. In the second part of the assignment you will use the constructed features for face recognition using your favorite classifier. In the third part of this assignment you will participate in a kaggle competition and attempt to improve your classification pipeline in any way you please. Lastly, you will reflect on the entire assignment and write a final discussion, this is on top of the discussion and comments that you already provided in the previous parts (detailing how and why you implemented certain methods). In summary, the assignment is broken down into four main parts:

1. build feature representations using handcrafted and non-handcrafted techniques (Sect. [2.1](#));
2. use and compare the feature representations in context of classification of faces (Sect. [2.2](#));
3. attempt to improve your model and participate in the kaggle competition (Sect. [2.3](#));

4. discussion (sect. [2.4](#)).

Your submission should be well documented, TAs should have to decipher your code to understand what you did. However, the goal is not to write a textbook, it is far more important to explain why you did things. Show your insights and reflection.

1.1 Preparation: learning to use a machine learning library

If you choose to build a (deep) learning network in section [2.2](#) or [2.3](#), we recommend using [Pytorch](#) or [TensorFlow/Keras](#)¹. Pytorch is very flexible, but it will require you to write more code, while Keras is a high-level neural networks API with prepared functions for standard training. For Pytorch, Pytorch Lightning would have equivalent functionality to Keras.

1.2 A faces dataset

In this assignment we use a small subset of the VGG Face Dataset [[PVZ15](#)]. The subset contains images of four celebrities, more specifically:

- 30 images of Jesse Eisenberg;
- 30 images of Mila Kunis;
- 10 images of Michael Cera (who arguably looks similar to Jesse Eisenberg);
- 10 images of Sarah Hyland (who arguably looks similar to Mila Kunis).

Use OpenCV to detect the faces in the images and extract the faces (this code is already provided in the template notebook but feel free to improve on it).

Hint: split this dataset (hold-out or cross-validation), being methodologically correct is in your best interest!

2 Assignment

2.1 Build feature representations

Once you know what defines an object or what allows you to differentiate between two objects, it becomes much easier to detect an object in an image. Learning a robust (same features are extracted from the same object in different conditions) and discriminative (different image objects can be easily separated from each other in feature space) feature representation is key to perform automated object detection effectively. Two big classes of feature representation methods can be identified:

- Handcrafted features (Sect. [2.1.1](#))
- Features learned from data (Sect. [2.1.2](#))

You will try to build both types of features on the faces dataset and use them as baseline.

¹Familiarize yourself with TensorFlow/Keras using for instance [this tutorial](#), created by François Chollet.

2.1.1 Handcrafted features

Handcrafted feature descriptors are extracted from the image itself, using a specific deterministic algorithm that might differ from task to task. To assess if a handcrafted feature is discriminative we compute it for a few images, we then compare the match score (e.g. euclidean distance between the feature representations) and observe if this matches our intuition/domain knowledge. Ideally, a (local) feature descriptor holds the robust and discriminative property and will additionally be invariant to e.g. illumination, pose, scale. This allows you to recognise objects of interest in virtually any image using that descriptor. Handcrafted features were commonly used in traditional machine learning approaches but we still see them today in complicated machine learning tasks, e.g. doing machine learning on a manifold mesh.

Compute features from pixel intensities in order to learn a representation of local structures present in the images. You can choose between one of these two state-of-the-art handcrafted image feature descriptors (feel free to use a library for this, e.g. OpenCV):

- Histogram of Oriented Gradients (more details can be found [here](#))
- Scale Invariant Feature Transform (see course book SB 4.1.2)

Carefully choose the hyperparameters. Keep in mind: how local should the descriptors be for the object that you want to detect; how to make your descriptor behave well in different circumstances (e.g. image with different lighting)? Visualize your feature representations to gain more insight. To do this, use the high dimensional data visualisation tool t-SNE² [MH08]), an example is shown in Fig. 1. Reflect on the parameters you've chosen. How does this feature compare to your previous grabbing task in the individual assignment? Did you need specific pre-processing steps before computing these feature descriptors on your images (which ones and why)? Did the visualisation show good discriminative and robustness properties?

2.1.2 Learning features from data: PCA

An alternative approach to handcrafting features that capture desired object properties is to learn these features from training data. In this case, the algorithm looks for patterns in the data that are informative.

Apply Principal Components Analysis (PCA) to find the different components of variation present in your faces training set and use the scores to form a feature vector. Think about this problem mathematically:

- How to convert an image (you can work in color if you like) dataset to a 2D matrix?
- Can you exploit the dimensionality of this data matrix to make your computations more effective?
- Mean subtraction or not?
- Shall we use eigenvalue or singular value decomposition?

²Scikit-learn has an [implementation](#).

- How many non-zero eigenvalues/singular values should we have?

Attempt to choose an *optimal* number p of principal components such that the dimensionality of the feature space is reduced but still informative. Project all images onto these same p -components and examine the effect on the space of faces in the dataset. Visualize the reconstructions of one face using gradually more eigenfaces and visualize the faces on the feature space using the first two principal components, an example is shown in Fig. 2. Reflect on your choice of p and how this might have influenced the performance of the features that you learned. Did you need specific pre-processing steps before computing these feature descriptors on your images (which ones and why)? How many non-zero eigenvalues did you have, why is this? Did the visualisation show good discriminative and robustness properties?

2.2 Classification

Now it's time to use the feature representations that you have constructed and start to classify images. In this exercise you will train a classifier that can recognise person Jesse Eisenberg (class 1) and Mila Kunis (class 2) or recognise that neither of them are present in the image (class 3). This can be achieved using 2 binary-classifiers (a Jesse recogniser and a Mila recogniser) or you can use one multi-classifier, the choice is yours! Feel free to use your favorite classification algorithm, you can use existing functions from open source Python libraries (e.g. scikit-learn) or implement them yourself. Evaluate the features built in Sect. 2.1 and report the mean classification accuracy on Jesse/Mila test images. How did you build this classifier for each of your feature representations? Do the models hold up when you predict on images of Michael Cera and Sarah Hyland?

Hint: Don't forget to apply the same pre-processing pipeline for each feature representation as you did before!

2.3 Improve performance

Now that you are familiar with the basics of computer vision and classification scenarios and have dipped your feet in traditional techniques it is time to experiment with the state of the art.

Improve on the parts of your pipeline that will have the biggest impact on your performance (face detection, feature extraction and/or classification model) and participate in the Kaggle competition. There are no rules for this section, you can try anything you please, just keep in mind that your data and compute are limited. Clearly document your zero-to-hero journey in the submitted notebook: Which iterations did you go through? Do you notice improvements?

Potentially you will run a large number of experiments, requiring too much time and memory for the Kaggle submission. If that's the case, simply document the earlier experiments and run only the final (best) pipeline. Please ensure that the TAs get a clear understanding of what you are doing from the notebook as this will be our basis for evaluation your work: no documentation $\Rightarrow$ lower grade.

Note: This competition is still in educational context, therefore final performance achieved in the Kaggle competition only contributes a small part of the grades earned in this assignment. A much larger part depends on your ingenuity, methodological correctness and insight. The Kaggle ranks are mainly intended for you to track the progress of your improvements and get an idea of what results others achieve on this task.

2.4 Discussion

Summarize (very briefly) what you have learned in this assignment. Discuss qualitative differences between the different feature representations that you constructed. Why/when would one work better than the other? Why is performance (sub) optimal? What would you do better/more if you would had plenty of time to spend on each step?

3 Practicalities

Make use of the *Group assignment 1* forum/discussion board on Toledo if you have any questions. Assignment submission and participation are performed through Kaggle.

3.1 Kaggle: Joining the competition

Closely follow [these instructions](#) for joining the competition, sharing your notebook with the TAs and making a valid notebook submission to the competition. You can only join the competition via this link. New Kaggle users may need to verify their account using their phone number; you may not be able to join the competition otherwise.

3.2 Submission

The notebook you submit for grading is the last notebook you submit in the [Kaggle competition](#) prior to the deadline (you can submit as many times as you want, we strongly recommend that you do so regularly). Only submissions through Kaggle will be accepted (time and memory issues are your own responsibility). Any submissions after the deadline will be disregarded.

A notebook submission not only lets you participate in the competition, it also runs the entire notebook and saves its output as if it were a report. This way it becomes an all-in-one-place document for the TAs to review. As such, please make sure that your final submission notebook is self-contained and fully documented. Provide strong arguments for the design choices that you made and what insights you got from your experiments. Most likely, this notebook format is not appropriate to run all your experiments at submission time (e.g. the training is a memory hungry and time consuming process) due to limited Kaggle resources. It can be a good idea to distribute your code otherwise and only summarize your findings, together with your final predictions, in the submission notebook. For example, you can substitute experiments with some text and figures that you have produced “offline”. We advise you to go through this document entirely before you really start. It can be a

good idea to then go through the template notebook and use this one as your first notebook submission to the competition.

To add your TAs to your Kaggle team, look for their Kaggle usernames on Toledo under *Didactic Team*.

References

- [MH08] Laurens van der Maaten and Geoffrey Hinton. “Visualizing data using t-SNE”. In: *Journal of machine learning research* 9.Nov (2008), pp. 2579–2605.
- [PVZ15] Omkar M. Parkhi, Andrea Vedaldi, and Andrew Zisserman. “Deep Face Recognition”. In: Section 3 (2015), pp. 41.1–41.12. DOI: [10.5244/c.29.41](https://doi.org/10.5244/c.29.41).



Figure 1: T-SNE plot (example from <https://towardsdatascience.com/t-sne-python-example-1ded9953f26>).




Figure 2: The face-feature plot (example from <https://sandipanweb.wordpress.com/2018/01/06/eigenfaces-and-a-simple-face-detector-with-pca-svd-in-python/>).